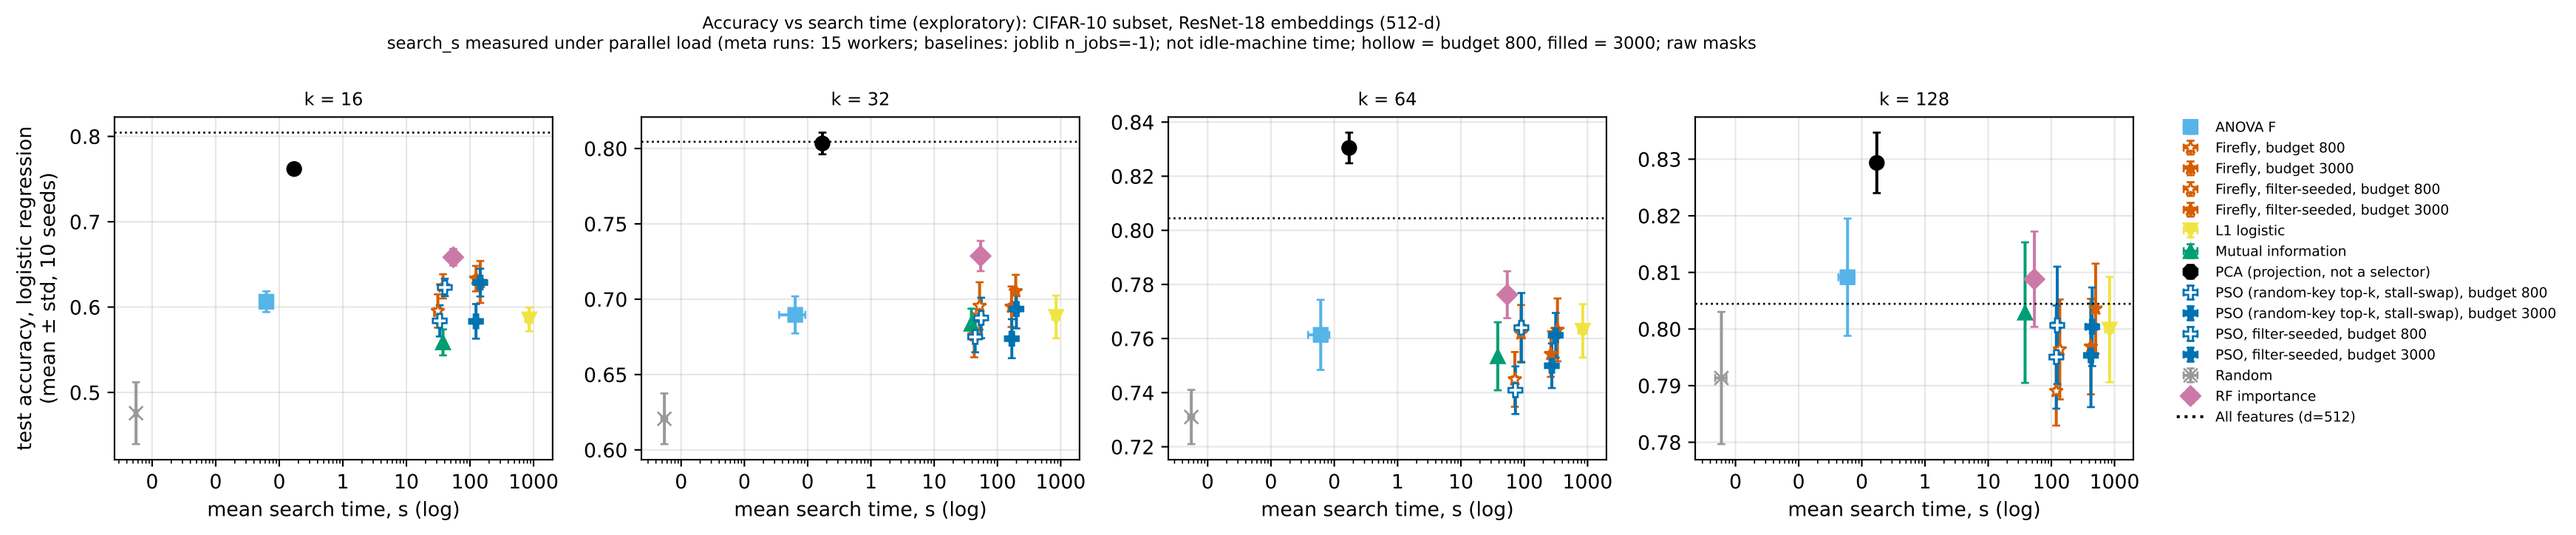
Accuracy vs search time (exploratory): CIFAR-10 subset, ResNet-18 embeddings (512-d)
search_s measured under parallel load (meta runs: 15 workers; baselines: joblib n_jobs=-1); not idle-machine time; hollow = budget 800, filled = 3000; raw masks
k = 16
k = 32
k = 64
k = 128
test accuracy, logistic regression
(mean ± std, 10 seeds)
mean search time, s (log)
ANOVA F
Firefly, budget 800
Firefly, budget 3000
Firefly, filter-seeded, budget 800
Firefly, filter-seeded, budget 3000
L1 logistic
Mutual information
PCA (projection, not a selector)
PSO (random-key top-k, stall-swap), budget 800
PSO (random-key top-k, stall-swap), budget 3000
PSO, filter-seeded, budget 800
PSO, filter-seeded, budget 3000
Random
RF importance
All features (d=512)